

In [1]:

```
1
2 import numpy as np # Linear algebra
3 import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
4 import plotly.offline as py
5 import plotly.graph_objs as go
6 import plotly.tools as tls
7 import plotly.express as px
8
9 from sklearn.feature_extraction.text import CountVectorizer
10 from keras.preprocessing.text import Tokenizer
11 from keras.preprocessing.sequence import pad_sequences
12 from keras.models import Sequential
13 from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
14 from sklearn.model_selection import train_test_split
15 from keras.utils.np_utils import to_categorical
16 import re
17
```

In [2]:

```
1 mydata = pd.read_csv('Sentiment.csv')
2 mydata = mydata[['text', 'sentiment']] #Selecting the required columns
```

In [3]: 1 mydata.head(15)

Out[3]:

	text	sentiment
0	RT @NancyLeeGrah: How did everyone feel about...	Neutral
1	RT @ScottWalker: Didn't catch the full #GOPdeb...	Positive
2	RT @TJMShow: No mention of Tamir Rice and the ...	Neutral
3	RT @RobGeorge: That Carly Fiorina is trending ...	Positive
4	RT @DanScavino: #GOPDebate w/ @realDonaldTrump...	Positive
5	RT @GregAbbott_TX: @TedCruz: "On my first day ...	Positive
6	RT @warriorwoman91: I liked her and was happy ...	Negative
7	Going on #MSNBC Live with @ThomasARoberts arou...	Neutral
8	Deer in the headlights RT @lizzwinstead: Ben C...	Negative
9	RT @NancyOsborne180: Last night's debate prove...	Negative
10	@JGreenDC @realDonaldTrump In all fairness #Bi...	Negative
11	RT @WayneDupreeShow: Just woke up to tweet thi...	Positive
12	Me reading my family's comments about how grea...	Negative
13	RT @ArcticFox2016: RT @AllenWestRepub "Dear @J...	Neutral
14	RT @pattonoswalt: I loved Scott Walker as Mark...	Positive

#1) Print the total number of positive and negative sentiments.

After observing the above dataset, Let's drop the 'Neutral' sentiments as the objective is to only calculate positive and negative tweets. Then let's filter the tweets so that only valid texts and words remain. Finally let's define the number of max features as 2000 and use Tokenizer to vectorize and convert text into Sequences so that the Network can deal with it as input.

```
In [4]: 1 mydata = mydata[mydata.sentiment != "Neutral"]
2 mydata['text'] = mydata['text'].apply(lambda x: x.lower())
3 mydata['text'] = mydata['text'].apply((lambda x: re.sub('[^a-zA-z0-9\s]', '', x)))
4
5 print(mydata[ mydata['sentiment'] == 'Positive'].size)
6 print(mydata[ mydata['sentiment'] == 'Negative'].size)
7
8 for idx,row in mydata.iterrows():
9     row[0] = row[0].replace('rt', ' ')
10
11 max_fatures = 2000
12 tokenizer = Tokenizer(num_words=max_fatures, split=' ')
13 tokenizer.fit_on_texts(mydata['text'].values)
14 X = tokenizer.texts_to_sequences(mydata['text'].values)
15 X = pad_sequences(X)
```

4472

16986

```
In [5]: 1 fig = px.histogram(mydata, x="sentiment")
2 fig.update_traces(marker_color="turquoise",marker_line_color='rgb(8,48,107)',
3                   marker_line_width=1.5)
4 fig.update_layout(title_text='Analyzing Different Sentiments')
5 fig.show()
```

The above plot shows the count of positive and negative sentiments available in the dataset.

#2) Build a sequential LSTM model to predict positive and negative sentiments.

```
In [6]: 1 embed_dim = 128
2 lstm_out = 196
3
4 model = Sequential()
5 model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
6 model.add(SpatialDropout1D(0.4))
7 model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
8 model.add(Dense(2,activation='softmax'))
9 model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy'])
10 print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 28, 128)	256000

spatial_dropout1d (SpatialDr	(None, 28, 128)	0

lstm (LSTM)	(None, 196)	254800

dense (Dense)	(None, 2)	394
=====		
Total params: 511,194		
Trainable params: 511,194		
Non-trainable params: 0		

None		

Hereby I declare the train and test dataset.

```
In [8]: 1 Y = pd.get_dummies(mydata['sentiment']).values
2 X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size = 0.3, random_state = 40)
3 print(X_train.shape,Y_train.shape)
4 print(X_test.shape,Y_test.shape)
```

```
(7510, 28) (7510, 2)
(3219, 28) (3219, 2)
```

Now, we train the model. Let's run 10 epoch

```
In [10]: 1 batch_size = 32
          2 model.fit(X_train, Y_train, epochs = 10, batch_size=batch_size, verbose = 2)
```

```
Epoch 1/10
235/235 - 35s - loss: 0.4201 - accuracy: 0.8206
Epoch 2/10
235/235 - 35s - loss: 0.3169 - accuracy: 0.8643
Epoch 3/10
235/235 - 35s - loss: 0.2733 - accuracy: 0.8846
Epoch 4/10
235/235 - 35s - loss: 0.2439 - accuracy: 0.8968
Epoch 5/10
235/235 - 35s - loss: 0.2157 - accuracy: 0.9067
Epoch 6/10
235/235 - 35s - loss: 0.1913 - accuracy: 0.9210
Epoch 7/10
235/235 - 35s - loss: 0.1825 - accuracy: 0.9244
Epoch 8/10
235/235 - 35s - loss: 0.1592 - accuracy: 0.9318
Epoch 9/10
235/235 - 35s - loss: 0.1485 - accuracy: 0.9366
Epoch 10/10
235/235 - 35s - loss: 0.1329 - accuracy: 0.9422
```

```
Out[10]: <keras.callbacks.History at 0x7fc17772c310>
```

Validation set extraction

Also measuring score and accuracy.

```
In [11]: 1 validation_size = 1500
2
3 X_validate = X_test[-validation_size:]
4 Y_validate = Y_test[-validation_size:]
5 X_test = X_test[:-validation_size]
6 Y_test = Y_test[:-validation_size]
7 score,acc = model.evaluate(X_test, Y_test, verbose = 2, batch_size = batch_size)
8 print("score: %.2f" % (score))
9 print("acc: %.2f" % (acc))
```

```
54/54 - 2s - loss: 0.5102 - accuracy: 0.8482
score: 0.51
acc: 0.85
```

Finally measuring the number of correct guesses. It is clear that finding negative tweets goes very well for the Network but deciding whether is positive is not really. My educated guess here is that the positive training set is dramatically smaller than the negative, hence the "bad" results for positive tweets.

```

In [12]: 1 pos_cnt, neg_cnt, pos_correct, neg_correct = 0, 0, 0, 0
          2 for x in range(len(X_validate)):
          3
          4     result = model.predict(X_validate[x].reshape(1,X_test.shape[1]),batch_size=1,verbose = 2)[0]
          5
          6     if np.argmax(result) == np.argmax(Y_validate[x]):
          7         if np.argmax(Y_validate[x]) == 0:
          8             neg_correct += 1
          9         else:
          10             pos_correct += 1
          11
          12     if np.argmax(Y_validate[x]) == 0:
          13         neg_cnt += 1
          14     else:
          15         pos_cnt += 1
          16
          17
          18
          19 print("pos_acc", pos_correct/pos_cnt*100, "%")
          20 print("neg_acc", neg_correct/neg_cnt*100, "%")
          21

```

```

1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s
1/1 - 0s

```


3) Based on the model, check the sentiment for the following two sentences

a. 'He is a great leader.'

b. 'He is a terrible leader.'

```
In [16]: 1 check_senti = ['He is a great leader.']
2 #vectorizing the tweet by the pre-fitted tokenizer instance
3 check_senti = tokenizer.texts_to_sequences(check_senti)
4 #padding the tweet to have exactly the same shape as `embedding_2` input
5 check_senti = pad_sequences(check_senti, maxlen=28, dtype='int32', value=0)
6 print(check_senti)
7 sentiment = model.predict(check_senti, batch_size=1, verbose = 2)[0]
8 if(np.argmax(sentiment) == 0):
9     print("negative")
10 elif (np.argmax(sentiment) == 1):
11     print("positive")
```

```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0 32  5  7 146 340]]
1/1 - 0s
positive
```

```
In [17]: 1 check_senti1 = ['He is a terrible leader.']
2         #vectorizing the tweet by the pre-fitted tokenizer instance
3         check_senti1 = tokenizer.texts_to_sequences(check_senti1)
4         #padding the tweet to have exactly the same shape as `embedding_2` input
5         check_senti1 = pad_sequences(check_senti1, maxlen=28, dtype='int32', value=0)
6         print(check_senti1)
7         sentiment = model.predict(check_senti1, batch_size=1, verbose = 2)[0]
8         if(np.argmax(sentiment) == 0):
9             print("negative")
10        elif (np.argmax(sentiment) == 1):
11            print("positive")
```

```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  32  5  7 1003 340]]
```

1/1 - 0s

negative